

LangChain and Retrieval Augmented Generation

LangChain is a framework for building applications that use large language models.

Retrieval Augmented Generation, commonly called RAG, combines information retrieval with language generation.

A RAG system first retrieves relevant documents or document chunks from a vector database.

The retrieved information is then provided to a language model as context.

ChromaDB is a vector database that can store embeddings and perform similarity searches.

Embeddings convert text into numerical vectors that capture semantic meaning.

Ollama can run language models locally.
This makes it possible to build local RAG applications.